TONER-1

260,032 parameters
int8 weights, one scale per row
Inference runs inside the PostScript interpreter

State of the art among language models that execute inside a print job.

SAMPLE OUTPUT

Unedited model output. No prompt: generation starts at the BOS token.

Once upon a time, there was a little girl named Lucy. She lived in a small house. One day, Lucy saw a sunny picture on the window. The sun was sad and rust slept. Luckie had a light bulb.

Then, a little boy named Tim came to the printer. He saw a long stack of paper on the wall and said, “Print, let’s hold something!” Lulu smiled and said, “Yes, I can come too!”

Tim showed a lot to someone else. Then he…

TOKEN STREAM

All 200 sampled tokens, one box per vocabulary piece. · marks a space, ¶ a newline.

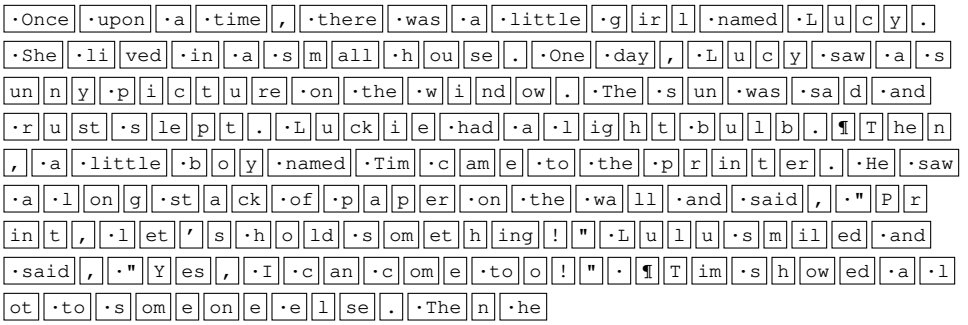

HOW THIS PAGE WAS MADE

1 HOST
Sends the print job. Performs no arithmetic.

→

2 INTERPRETER
Parses the int8 weights, then runs the forward pass once per token.

→

3 PAGE
The same job typesets the story and calls showpage.

TOKENS	SECONDS	TOKENS / SEC	SEED	TEMPERATURE
200	2.4	82.3	20	0.8

INTERPRETER	GPL Ghostscript · revision 10071 · PostScript version 3010
MODEL	stories260K printer-tuned (Llama-2 architecture) · dim 64 · 5 layers · 8 query / 4 KV heads · 512-token vocab
STOPPED BY	Token budget: STEPS = 200. The story was cut off mid-thought.

Host compute: 0 FLOPs. Inference cost includes toner.

Every word above was generated on-device, one token at a time, by the interpreter named in this footer; the host computer only sent bytes. Seconds are interpreter usertime spent generating. rand is implementation-defined, so the same seed on a different interpreter tells a different story.