TONER-1

260,032 parameters
int8 weights, one scale per row
Inference runs inside the PostScript interpreter

State of the art among language models that execute inside a print job.

SAMPLE OUTPUT

Unedited model output. No prompt: generation starts at the BOS token.

Once upon a time, there was a little girl named Lily. She loved to play outside and seek. One day, she went to the park with her mommy. They saw a cute dog and a small dog named Max who was playing with a ball. Lily was very excited and wanted to help her mommy.

Lily opened the dog and saw a big bell with a ball. She asked her mommy, “Can I play with you?” Her mommy said, “Okay, Max.” Lily was happy to hear that and said, “Thank you, Max. I am you so restless.”

TOKEN STREAM

All 184 sampled tokens, one box per vocabulary piece. · marks a space, ¶ a newline.

`·Once` `·upon` `·a` `·time` `,` `·there` `·was` `·a` `·little` `·g` `ir` `l` `·named` `·Lily` `.` `·She`
`·lo` `ved` `·to` `·play` `·` `out` `s` `id` `e` `·and` `·s` `e` `e` `k` `.` `·One` `·day` `,` `·she` `·w` `ent`
`·to` `·the` `·p` `ar` `k` `·with` `·her` `·mom` `my` `.` `·They` `·saw` `·a` `·c` `ut` `e` `·do` `g` `·and`
`·a` `·s` `m` `all` `·do` `g` `·named` `·M` `a` `x` `·w` `h` `o` `·was` `·play` `ing` `·with` `·a` `·b` `all` `.`
`·Lily` `·was` `·very` `·e` `x` `c` `it` `ed` `·and` `·want` `ed` `·to` `·he` `l` `p` `·her` `·mom` `my` `.` `¶`
`L` `ily` `·o` `p` `en` `ed` `·the` `·do` `g` `·and` `·saw` `·a` `·big` `·be` `ll` `·with` `·a` `·b` `all` `.`
`·She` `·a` `s` `ked` `·her` `·mom` `my` `,` `·"` `C` `an` `·I` `·play` `·with` `·you` `?` `"` `·H` `er` `·mom`
`my` `·said` `,` `·"` `O` `k` `ay` `,` `·M` `a` `x` `.` `"` `·Lily` `·was` `·happy` `·to` `·he` `ar` `·that`
`·and` `·said` `,` `·"` `T` `h` `an` `k` `·you` `,` `·M` `a` `x` `.` `·I` `·a` `m` `·you` `·so` `·` `re` `st` `le` `s`
`s` `.` `"`

HOW THIS PAGE WAS MADE

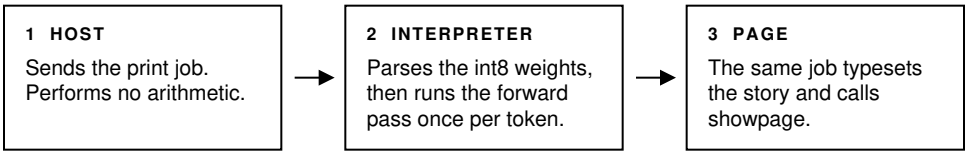

TOKENS	SECONDS	TOKENS / SEC	SEED	TEMPERATURE
184	3.5	52.3	20	0.8

INTERPRETER	GPL Ghostscript · revision 10071 · PostScript version 3010
MODEL	stories260K (Llama-2 architecture) · dim 64 · 5 layers · 8 query / 4 KV heads · 512-token vocab
STOPPED BY	End-of-story token, after 184 tokens

Host compute: 0 FLOPs. Inference cost includes toner.

Every word above was generated on-device, one token at a time, by the interpreter named in this footer; the host computer only sent bytes. Seconds are interpreter usertime spent generating. rand is implementation-defined, so the same seed on a different interpreter tells a different story.