TONER-1

260,032 parameters
int8 weights, one scale per row
Inference runs inside the PostScript interpreter

State of the art among language models that execute inside a print job.

SAMPLE OUTPUT

Unedited model output. No prompt: generation starts at the BOS token.

One day, a white, wore a terrible pain on TV. Tom had some job to move before lunch. One day, the'ving girl applauded under the pain. She looked out the sea harbor and saw a small pit. The warm sky on the bulb and learned to be more cool.

Tom learned he's not more job because his bups lived near her futherck. But today, Jen stepped smaller, and showed day a ryster. It really want…

TOKEN STREAM

All 200 sampled tokens, one box per vocabulary piece. · marks a space, ¶ a newline.

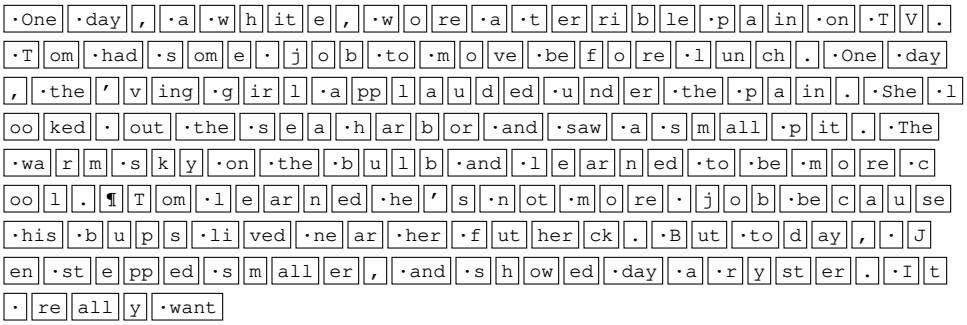

HOW THIS PAGE WAS MADE

1 HOST
Sends the print job. Performs no arithmetic.

→

2 INTERPRETER
Parses the int8 weights, then runs the forward pass once per token.

→

3 PAGE
The same job typesets the story and calls showpage.

TOKENS	SECONDS	TOKENS / SEC	SEED	TEMPERATURE
200	3.7	54.1	6	1.4

INTERPRETER	GPL Ghostscript · revision 10071 · PostScript version 3010
MODEL	stories260K (Llama-2 architecture) · dim 64 · 5 layers · 8 query / 4 KV heads · 512-token vocab
STOPPED BY	Token budget: STEPS = 200. The story was cut off mid-thought.

Host compute: 0 FLOPs. Inference cost includes toner.

Every word above was generated on-device, one token at a time, by the interpreter named in this footer; the host computer only sent bytes. Seconds are interpreter usertime spent generating. rand is implementation-defined, so the same seed on a different interpreter tells a different story.